DO

Document Organizer

COPY-ONLY DOCUMENT RE-ORGANIZER

RELEASE NOTES V2.1

Release Notes v2.1

What is new, what changed, verified results and known limitations

Document ID DO-RN-07 · **Version** 2.1 · **Date** 22-Oct-2026

Prepared by Yasin Jariwala

Audience Everyone installing, upgrading or evaluating Document Organizer

Document control

Document	Release Notes v2.1
Document ID	DO-RN-07
Version	2.1
Date	22-Oct-2026
Prepared by	Yasin Jariwala
Product	Document Organizer v2.1
Audience	Everyone installing, upgrading or evaluating Document Organizer
Classification	Public
Source	github.com/yasin679/Document_Organizer · docs/release-notes.md

Version history

Version	Date	Author	Changes
1.0	06-Oct-2026	Yasin Jariwala	First release
2.0	13-Oct-2026	Yasin Jariwala	General-purpose release: Amazon S3 and S3-compatible storage, copy between providers, container hosting; tenant ID field removed
2.1	22-Oct-2026	Yasin Jariwala	Several mapping files per job; File name cells listing several files split into one entry per file; NULL placeholders treated as blank; repeats of missing records reported as MISSING

Related documents

ID	Document
DO-UG-01	User Guide
DO-IG-02	Installation and Administration Guide
DO-QS-03	Quick Start
DO-FAQ-04	FAQ and Troubleshooting
DO-TD-05	Technical Design
DO-UAT-06	UAT Test Cases

Contents

1 Release summary 3
2 New and fixed in v2.1 3
2.1 Several file names in one cell (fix) 3
2.2 Placeholder values (fix) 3
2.3 Repeated records of missing files (fix) 4
2.4 Several mapping files per job (new) 4
2.5 Excluding exceptions in bulk (new) 4
2.6 Rule for a file that matches several records (new) 5
2.7 Other changes 5
3 Verified results of v2.1 5
3.1 Replay of a real job 5
3.2 Automated tests 6
4 Upgrading from v2.0 7
5 New in v2.0 7
5.1 Amazon S3 and S3-compatible storage 7
5.2 Copy between providers 7
5.3 Hosting and packaging 8
5.4 Configuration 8
5.5 User interface 8
5.6 Documentation 8
6 Fixed in v2.0 8
7 Verified results of v2.0 9
8 Changes that affect upgrades from v1.0 9
9 Known issues and limitations 10
10 Installation notes 10
11 Roadmap 10

1 Release summary

Item	Value
Product	Document Organizer
Version	2.1
Release date	22-Oct-2026
Previous version	2.0 (13-Oct-2026), notes from section 5
Status	Ready for UAT
Source	https://github.com/yasin679/Document_Organizer
Platform	.NET 8 on Windows, Linux or macOS; Docker image for any container platform

Document Organizer copies employee documents into a clean, consistent structure, `EmployeeCode / DocumentType / Year / FileName`, using mapping files to identify each file. It works with **Azure Blob Storage**, **Amazon S3** (and S3-compatible storage) and **local or network folders**, in any combination. It is copy-only: source documents are never moved, changed or deleted.

Version 2.1 fixes how MISSING and DUPLICATE are counted for real-world mapping files, adds **several mapping files per job**, and lets operators **exclude exceptions in bulk** with a note.

2 New and fixed in v2.1

2.1 Several file names in one cell (fix)

Some document logs list several uploaded files for one document in a single File name cell, for example `marksheet_17206.jpg,marksheet_17206.pdf`. Earlier versions looked for one file with that whole text as its name, so the record was reported MISSING and the files themselves UNMAPPED.

- Each name now becomes its own entry with its own status and destination. Record `ROW-68` becomes `ROW-68/1`, `ROW-68/2` and so on.
- Commas, semicolons and line breaks separate names; spaces around them are ignored.
- A name that is exactly the name of a file in the source (a real file name containing a comma) is kept whole, also when it sits next to other names in the cell (`report, final.pdf,x.pdf`). Folder paths in the cell work.
- A cell that only holds placeholders (`NULL,NULL`) is blank, so INVALID.
- The option **Several file names in one cell** in step 4 is on by default.

2.2 Placeholder values (fix)

Cells that only contain `NULL`, `N/A`, `NA`, `#N/A`, `None`, `Nil`, `-`, `--` or `(null)` (typical of database exports) are now treated as empty. A File name of `NULL` is therefore INVALID ("Blank FileName") instead of MISSING.

2.3 Repeated records of missing files (fix)

A repeated mapping record (same file, employee and type as an earlier record) was always reported DUPLICATE, even when the file did not exist and the earlier record was MISSING. Now:

- DUPLICATE when the earlier record found its file (the file is copied once).
- MISSING when the file is not in the source; the reason names the earlier record.

The DUPLICATE count now means "copied once already", and the MISSING count covers every record that needs follow-up.

2.4 Several mapping files per job (new)

- Step 2 accepts **one or more files** at once, and more can be added or removed until the job is approved.
- Records of all files are combined in one scan; a document repeated across files is reported once as DUPLICATE, naming the other file.
- Each file has **its own column mapping** in step 3, so exports with different column names, column order, extra columns, title rows, CSV delimiters and file types (`.xlsx` and `.csv`) can be combined. The scan is refused, naming the file, until every file has its required columns.
- A file without the year rule's column uses its other date column: a Year-only file under the **UploadedOn year** rule uses its Year column, and the reverse.
- Upload checks: the same content twice is refused; two files with the same name become `name (2).csv`; a failed file cancels the whole upload without leftovers.
- Every record keeps the **name of its mapping file**: shown in the results when a job has several files, in a new **Mapping File** column of both reports, and in the report summary (each file with rows and SHA-256).
- New folder-template token `{MappingFile}` (file name without extension) keeps each file's documents apart.
- New API endpoints: `POST /jobs/{id}/mappings` (add files) and `DELETE /jobs/{id}/mappings/{mappingId}`; `POST /jobs/{id}/mapping` now replaces all files with the uploaded file(s).

2.5 Excluding exceptions in bulk (new)

Exceptions that need no follow-up can be set aside in bulk, with a required note, instead of staying in the exception lists.

- Works for every exception type: unmapped, missing, ambiguous, duplicate, validation errors and failed.
- Three ways to choose: **tick rows** (with a select-page box), **everything in the current view** (tab and search), or a **whole reason** from the new *Exceptions by reason* list (for example 3,230 *Blank file name* at once).
- Excluded items get the new status **EXCLUDED**. They are never copied, no longer count as exceptions, and keep their original status and reason. Both reports have four new columns at the end: Excluded From Status, Excluded By, Excluded At (UTC), Exclusion Note.
- Exclusions are remembered: a new scan excludes the same file or record again while its status is unchanged.
- **Restore** from the Excluded tab in the same three ways.
- Allowed while a job is awaiting approval, paused or completed. Excluding the last exceptions of a completed job makes it **Completed**.

- Every exclusion and restore is in the audit log with the count, statuses, filter and note. Viewers can see but not change exclusions.
- On the real-job replay, excluding all 56,019 exceptions took 1.8 seconds, restoring them 1.5 seconds.
- New API: `GET /jobs/{id}/reasons`, `POST /jobs/{id}/items/exclude`, `POST /jobs/{id}/items/restore`; `GET /jobs/{id}/items` accepts `group`.

2.6 Rule for a file that matches several records (new)

Step 4 has a new rule, **When a file matches several mapping records**:

- **Flag as ambiguous** (default, as before): the file is AMBIGUOUS until a person picks the record.
- **Keep all**: the file is copied once for every record that names it, each to its own destination, and each copy is its own row in the validation and reconciliation reports (so the final Excel lists every record). If two of those records would give the same destination path, one copy is kept and the other is DUPLICATE. The source file is still counted once in *Files discovered*.

The rule is shown in the approval panel, the report summary and the audit entry for the rules.

2.7 Other changes

- Matching uses an internal sequence number, so repeated record IDs (within or across mapping files) no longer affect matching.
- Resolving an ambiguous file identifies the record by ID and mapping file.
- CSV files saved by Excel on Windows (Windows-1252) are read correctly; earlier, accented file names in them were not found.
- In an Excel workbook, the first visible sheet with a header row is used, so a cover or summary sheet in front is fine.
- Adding a second file to a job created in v2.0 keeps that job's hand-edited column mapping.
- Saving the rules without per-file columns (v2.0 clients and scripts) no longer resets other files' columns.
- Resolving an ambiguous file to a record whose ID also exists in another mapping file no longer marks it DUPLICATE by mistake.
- A browser icon was added (removes a 404 on every page load).
- New automated tests `tests/e2e-mapping.mjs` and `tests/e2e-scenarios.mjs`, both run in GitHub Actions.

3 Verified results of v2.1

3.1 Replay of a real job

A real job's reconciliation report (51,300 mapping records, 107,799 source files in a local folder) was rebuilt with placeholder file contents and scanned with both versions. Version 2.0 reproduced the original report exactly.

Status	v2.0	v2.1	Change explained
Processed	41,296	55,948	5,840 records listed 14,657 files in one cell; they are now split and matched

Status	v2.0	v2.1	Change explained
Unmapped	66,502	51,850	The same 14,652 files now have a record
Missing	8,351	923	Multi-file cells are no longer one "missing" file; 1,591 `NULL` file names are now INVALID; 3 repeats of missing records added
Invalid	1,639	3,231	1,592 `NULL` file names (blank)
Duplicate	14	15	4 repeats of missing records left DUPLICATE; 5 files listed in two records appeared after splitting
Ambiguous / Failed	0 / 0	0 / 0	

Every one of the 14,657 file names in multi-file cells exists in the source. Scan: 4.8 seconds; copy of 55,948 files: 7.8 seconds (local disk).

3.2 Automated tests

Test	Result
`tests/e2e-mapping.mjs`: split cells, real comma in a file name, `NULL`, repeats of missing and found records, two mapping files with different column names, remove and add, `{MappingFile}`, report column	24 checks pass
`tests/e2e-exclusions.mjs`: exclude ticked rows, a whole reason and a searched view; note required; ready items refused; roles; remembered after a new scan (and not when the status changed); restore; completed job becomes Completed; report columns; audit	28 checks pass
Browser test of bulk exclusion: tick boxes, select count, note required, exclude a reason, exclude a view, Excluded tab, restore all	14 checks pass, no browser errors
`tests/e2e-scenarios.mjs`: 20 kinds of File name cell; four mapping files with different columns, structure, encoding and date columns in one job; hand-mapping, per-file year rule, ambiguity across files with colliding IDs; upload checks, roles and locked jobs; workbook with a summary sheet; the same mapping through folder → Azure → S3 → folder; rule for files matching several records (scenario F)	87 checks pass
`tests/e2e.mjs`: 7 jobs across folders, Azure and S3, including cross-provider	159 checks pass, unchanged from v2.0
Upgrade: job with hand-edited columns made on v2.0, opened, rescanned and given a second file on v2.1	6 checks pass
Browser test: upload four files, tabs, hand-mapping, save, resolve, remove, same file twice	15 checks pass, no browser errors

4 Upgrading from v2.0

Item	What happens
Configuration	No changes needed
Database	Columns `jobs.mappings_json` and `items.mapping_file` are added automatically at start-up
Existing jobs	Keep their single mapping file and results. Scanning one again applies the new rules, so its counts can change as in section 3.1.
Reports	Both reports have a new **Mapping File** column after Mapping Record ID, and four exclusion columns at the end. Update anything that reads columns by position.
Database	Also adds the `exclusions` table and four `excluded_*` / `exclusion_note` columns to `items`
API	`POST /jobs/{id}/mapping` accepts several files and replaces all mapping files; use `POST /jobs/{id}/mappings` to add

5 New in v2.0

5.1 Amazon S3 and S3-compatible storage

- **Amazon S3** as a source and destination, alongside Azure Blob Storage and folders.
- Works with **S3-compatible services** (MinIO, Wasabi, Cloudflare R2 and others) through `ServiceUrl`, with path-style addressing.
- **Credentials:** access key, or the standard AWS chain: environment variables, EKS web identity (IRSA), `~/.aws/credentials` profile, ECS / EKS Pod Identity container role, EC2 instance role (IMDSv2).
- **Server-side copy** (CopyObject) between buckets on the same connection, with the source ETag as a condition.
- **Multipart upload** for files over 8 MiB; each part carries an MD5 check; failed uploads are aborted.
- **Conditional writes** (`If-None-Match: *`) so an existing object is never replaced unless Overwrite is allowed.
- Objects in **Glacier Flexible Retrieval** and **Deep Archive** are reported as INVALID (archive tier) and never restored automatically.
- A wrong bucket region is reported with the bucket's real region.
- The destination bucket is created on **Check access** if it does not exist and the credentials allow it.
- **Content-MD5** integrity checks use the S3 ETag where it equals the MD5 (single-part uploads without KMS or customer-key encryption).
- Implemented on the S3 REST API with Signature Version 4; no AWS SDK is required.

5.2 Copy between providers

- Source and destination can be **different storage types**: Azure to S3, S3 to Azure, either to or from a folder.
- Files are streamed, verified (SHA-256 or Content-MD5) and tagged with traceability metadata on the destination.

- Source metadata is carried across providers, with names adjusted to each provider's rules.

5.3 Hosting and packaging

- **Dockerfile:** small runtime image, non-root user, port 8080, data on a volume.
- **docker-compose.yml:** the application plus Azurite (Azure emulator) and MinIO (S3 emulator) for a trial without any cloud account.
- `/healthz` endpoint for container and load-balancer health checks.
- `BehindReverseProxy` setting to trust `X-Forwarded-For` / `X-Forwarded-Proto` behind HTTPS proxies and load balancers.
- `start.sh` for Linux and macOS, next to `start.cmd` for Windows.
- **GitHub Actions** workflow that builds the application and runs the end-to-end tests against Azurite and moto on every push.
- `.gitignore`, `.gitattributes`, `.env.example` and a sample mapping file (`samples/sample-mapping.csv`).

5.4 Configuration

- Storage connections are **typed** with `Provider: AzureBlob` or `Provider: S3` (inferred if omitted).
- Any connection can be marked `UseForDemo`; **Load demo data** then also uploads the sample files there.
- Secrets can live in `appsettings.Production.json`, which is excluded from Git, or in environment variables.

5.5 User interface

- **New job** offers three storage types. Only connections of the chosen type are listed, and the box is labelled Container or Bucket accordingly.
- **Settings** has a switch for Amazon S3 and shows each connection's type.
- The jobs table shows the **destination** of each job.
- Demo shortcuts are shown for every demo connection that was loaded.
- The rule "Keep blob metadata" is now **Keep object metadata** and applies to every cloud destination.

5.6 Documentation

All seven guides were rewritten for v2.0 and are kept in the repository under `docs/`, as Markdown (readable on GitHub) and as PDF.

6 Fixed in v2.0

#	Issue	Fix
1	Copies from Azure Blob Storage to a non-Azure destination (for example a local folder) failed every file with "Source file changed after the scan"	Azure's list and properties calls return the ETag in different forms (unquoted and quoted). ETags are now stored in one form and compared without quotes. Azure-to-Azure copies were not affected.

#	Issue	Fix
2	Source metadata copied to Azure could be rejected if a name contained characters Azure does not allow	Metadata names are adjusted to the destination's rules

7 Verified results of v2.0

Test	Result
End-to-end automated test: 7 jobs covering folder, Azure (Azurite), S3, Azure → folder, S3 → folder, Azure → S3, S3 → Azure	159 checks pass
S3 against an emulator that enforces Signature V4, about 10,000 requests	All accepted; wrong secret rejected with `SignatureDoesNotMatch`
20 MB file to S3; key with spaces, accents and `+&=@#`	Multipart in 3 parts, byte-identical; exact key preserved
Second run into the same destination	Every item DUPLICATE; nothing overwritten
66,000-row mapping file with 60,040 files (local folder)	Read, matched and validated in about 3 seconds
v1.0 scale test (176,000 rows, 60,000 files; restart during copy and resume)	Exactly 60,000 processed, no duplicates

8 Changes that affect upgrades from v1.0

Change	Effect	What to do
SignUpID removed	The SignUpID field on New job, the SignUpID column mapping, the `{SignUpId}` template token and the SignUpID report columns are gone. The mapping file no longer needs a tenant column.	Remove `{SignUpId}` from any saved templates. Extra columns in a mapping file are simply ignored.
Metadata keys renamed	Copied files are tagged `docorgjobid`, `docorgrecordid` and `docorgemployeecode`	Update anything that reads the old keys
Connection `Provider`	New optional setting; inferred when missing	Add `"Provider": "AzureBlob"` to existing connections for clarity
Product-neutral naming	No customer-specific branding, names or examples remain in the application or documents	None
Demo data	New fictional data set (employee codes starting with `E`); demo mapping file `demo-mapping.xlsx`	Load the demo data again
Default connections	`appsettings.json` now ships `Azurite` and `LocalS3` emulator connections	Remove them in production

The database is upgraded automatically on start; existing jobs, users and the audit log are kept.

9 Known issues and limitations

#	Item	Workaround or plan
1	Not yet tested with real Azure (Azure CLI, managed identity) or real AWS (IAM roles); automated tests use emulators	Run UAT cases DO-UAT-06 before using real documents
2	Local user accounts only	Single sign-on (OpenID Connect) on the roadmap
3	Archived objects (Azure Archive, S3 Glacier Flexible Retrieval and Deep Archive) are not restored	Restore them in the storage service, then scan again
4	Server-side copy only within one Azure storage account or one S3 connection	Other combinations stream through the application
5	One application instance per data folder	Do not run two instances against the same data volume
6	The Docker image has not been built in the release pipeline yet	Build it locally (`docker compose up --build`) and run UAT-01 and UAT-02
7	A File name cell is split on commas and semicolons unless the text is a source file name. A real file name with a comma that is **not** in the source is still split (then reported MISSING per part); so is a name with more than five commas	Turn off **Several file names in one cell** in step 4 for such a mapping file
8	The year rule is chosen for the whole job; only the fallback between Uploaded on and Year differs per file	Use one job per year rule if files need different rules (for example one by file name pattern)

10 Installation notes

- New installation: follow the Installation and Administration Guide (DO-IG-02).
- Requires the .NET 8 SDK (or ASP.NET Core 8 Runtime for a published build), or Docker.
- Set `EnableDemoData` to `false` and remove the emulator connections before production use.
- Keep secrets out of Git: environment variables or `appsettings.Production.json`.

11 Roadmap

Item	Planned
Sign-in	Single sign-on with OpenID Connect (Microsoft Entra ID, Okta, Google)
Storage	Google Cloud Storage; optional restore of archived objects with approval
Dashboard	Filters by job label, storage type and date
Packaging	Published container image and release builds on GitHub

Item	Planned
Intelligence	OCR and document classification, always with human approval