DO

Document Organizer

COPY-ONLY DOCUMENT RE-ORGANIZER

USER GUIDE

User Guide

Screen-by-screen guide for operators, viewers and administrators

Document ID DO-UG-01 · **Version** 2.1 · **Date** 22-Oct-2026

Prepared by Yasin Jariwala

Audience Operators, viewers and administrators of Document Organizer

Document control

Document	User Guide
Document ID	DO-UG-01
Version	2.1
Date	22-Oct-2026
Prepared by	Yasin Jariwala
Product	Document Organizer v2.1
Audience	Operators, viewers and administrators of Document Organizer
Classification	Public
Source	github.com/yasin679/Document_Organizer · docs/user-guide.md

Version history

Version	Date	Author	Changes
1.0	06-Oct-2026	Yasin Jariwala	First release
2.0	13-Oct-2026	Yasin Jariwala	General-purpose release: Amazon S3 and S3-compatible storage, copy between providers, container hosting; tenant ID field removed
2.1	22-Oct-2026	Yasin Jariwala	Several mapping files per job; File name cells listing several files split into one entry per file; NULL placeholders treated as blank; repeats of missing records reported as MISSING

Related documents

ID	Document
DO-IG-02	Installation and Administration Guide
DO-QS-03	Quick Start
DO-FAQ-04	FAQ and Troubleshooting
DO-TD-05	Technical Design
DO-UAT-06	UAT Test Cases
DO-RN-07	Release Notes v2.1

Contents

1 About this guide 3
1.1 What Document Organizer does 3
1.2 Who should read it 3
1.3 How to read the screenshots 3
1.4 Key terms 3
2 Getting started 5
2.1 Opening the application and signing in 5
2.2 Screen layout 5
2.3 Roles and what you can do 5
2.4 Signing out and session timeout 6
3 Dashboard 7
4 Jobs list 8
5 Creating a job: source and destination 9
5.1 Choose the source 9
5.2 Choose the destination and name the job 9
5.3 Filling in the demo locations 10
6 Uploading the mapping files 11
7 Mapping the columns 13
8 Setting the rules 14
9 Scanning and validating 15
10 Reviewing and approving 16
10.1 Results summary 16
10.2 Exceptions and results 16
10.3 Excluding exceptions in bulk 16
10.4 Resolving an ambiguous file 19
10.5 Checking the output and approving 19
10.6 Maker-checker (second approver) 20
11 Processing (copying) 23
12 Reconciliation and reports 25
12.1 What the reports contain 25
13 Audit log 26
14 Settings 27
14.1 Users and roles (administrators) 27
15 Status reference 29
16 Glossary 30
17 Appendix A: practise with demo data

1 About this guide

This guide walks through every screen of Document Organizer in the order you meet them when you run a job. Each screen has a screenshot with numbered markers, and a table that explains every marked item.

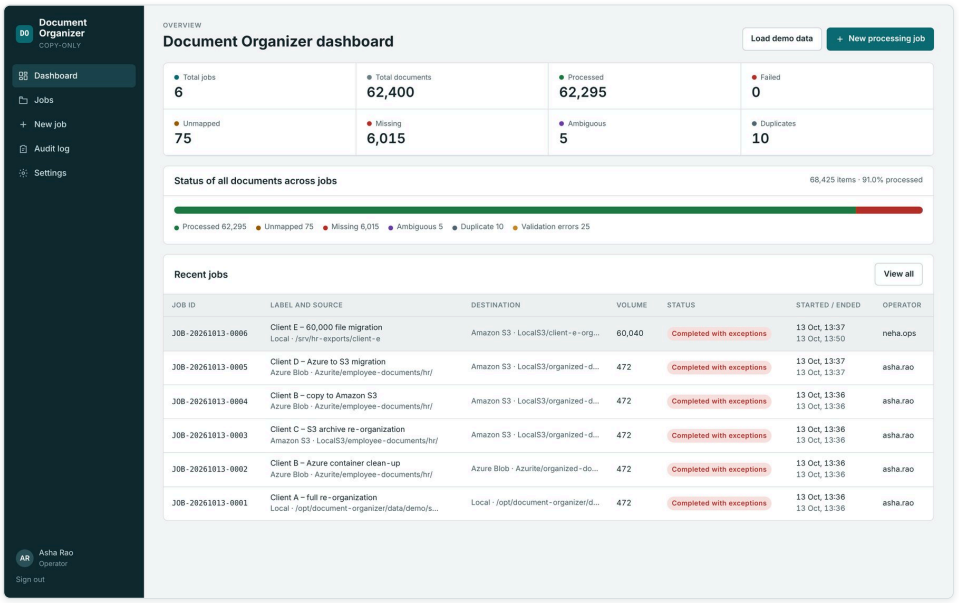


Figure 1: The Document Organizer dashboard after several completed jobs

1.1 What Document Organizer does

Document Organizer copies employee documents into a clean folder structure: `EmployeeCode / DocumentType / Year / FileName`. It uses a mapping file, usually the document log exported from the HR system where the documents were uploaded, to decide where each file belongs.

The documents can be in **Azure Blob Storage**, **Amazon S3** (or S3-compatible storage such as MinIO, Wasabi or Cloudflare R2), or a **local or network folder**. The organized copy can go to any of the three, so you can also use a job to move documents from one storage to another, for example from Azure to S3.

For example, the mapping record for employee E00879, document type *Medical Certificate*, uploaded on 25-Nov-2021, with file name `e00879_final_certificate.pdf`, is copied to:

```
E00879/Medical Certificate/2021/e00879_final_certificate.pdf
```

Three rules the application always follows

- **Copy only.** The original files are read, never moved, changed or deleted.
- **No guessing.** A file that cannot be placed with certainty is flagged for a person to decide.
- **Everything is accounted for.** Every file and every mapping record ends with a status and appears in the reports.

1.2 Who should read it

Reader	Sections to read
Operators (team members who run jobs)	All sections
Viewers (team leads, auditors, account managers)	2, 3, 4, 10, 12, 13 and 15
Administrators	All sections, plus the Installation and Administration Guide (DO-IG-02)

1.3 How to read the screenshots

Orange numbered markers on a screenshot point to the item described in the matching row of the table below it. The orange outline shows the exact area the marker refers to. Names of buttons and fields are shown in **bold**.

Sample data

The screenshots use the built-in demo data and a test job of 60,000 files. All employee codes, names, file names and client names in them are fictional.

1.4 Key terms

Term	Meaning
Job	One run that organizes one set of documents, from a source to a destination, with its own rules and reports.
Source	Where the documents are today: an Azure Blob container, an S3 bucket (each with an optional prefix) or a folder. Only read.

Term	Meaning
Destination	Where the organized copy is written. Must be separate from the source.
Mapping file	A spreadsheet (`.xlsx` or `.csv`) with one row per document. It says which employee and document type each file belongs to. A job can use several mapping files.
Container / bucket	The top-level storage area in Azure Blob Storage (container) or Amazon S3 (bucket).
Prefix	A "virtual folder" inside a container or bucket, for example `hr/` .
Storage connection	A named, pre-configured link to an Azure storage account or an S3 account. Set up by your administrator.
Status	The result for each file or mapping record, for example READY, UNMAPPED or PROCESSED. See section 15.

2 Getting started

2.1 Opening the application and signing in

Open the link your administrator gives you. On a laptop installation it is http://localhost:5080. The sign-in page opens.

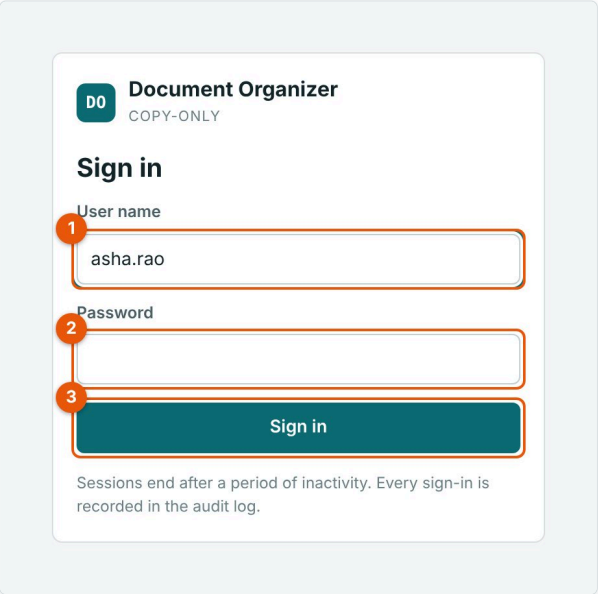


Figure 2: Sign-in page

#	Item	What to do
1	**User name**	Type the user name your administrator created for you.
2	**Password**	Type your password. Passwords are case-sensitive.
3	**Sign in**	Click to sign in. You land on the dashboard.

If sign-in fails

After 5 wrong passwords, sign-in from your computer is blocked for 10 minutes. Ask an administrator to reset your password if you have forgotten it. Every sign-in, successful or not, is recorded in the audit log.

2.2 Screen layout

Every page has the same layout: the navigation bar on the left, the page title at the top, and the page content below it.

Menu item	What it opens	Who sees it
Dashboard	Totals across all jobs and the most recent jobs	Everyone
Jobs	The full list of jobs	Everyone
New job	Starts a new processing job	Operators and administrators
Audit log	Every action taken in the application	Everyone
Settings	Settings, users and your password	Everyone (changes: administrators only)

2.3 Roles and what you can do

Action	Viewer	Operator	Administrator
View dashboard, jobs, results and audit log	Yes	Yes	Yes
Download validation and reconciliation reports	Yes	Yes	Yes
Create jobs, upload mapping files, set rules, scan	No	Yes	Yes
Resolve ambiguous files, approve, pause, resume, retry	No	Yes	Yes
Delete a job	No	Own draft jobs	Any job
Change settings, manage users	No	No	Yes
With maker-checker on: approve a job they did not work on, or a change request they did not make	No	Yes	Yes

2.4 Signing out and session timeout

Click **Sign out** at the bottom of the navigation bar. If you do nothing for 30 minutes (your administrator can change this) you are signed out automatically, and the next click takes you to the sign-in page. A job that is running keeps running after you sign out.

3 Dashboard

The dashboard gives an overview of all jobs: how many documents have been processed, how many need attention, and the latest jobs.

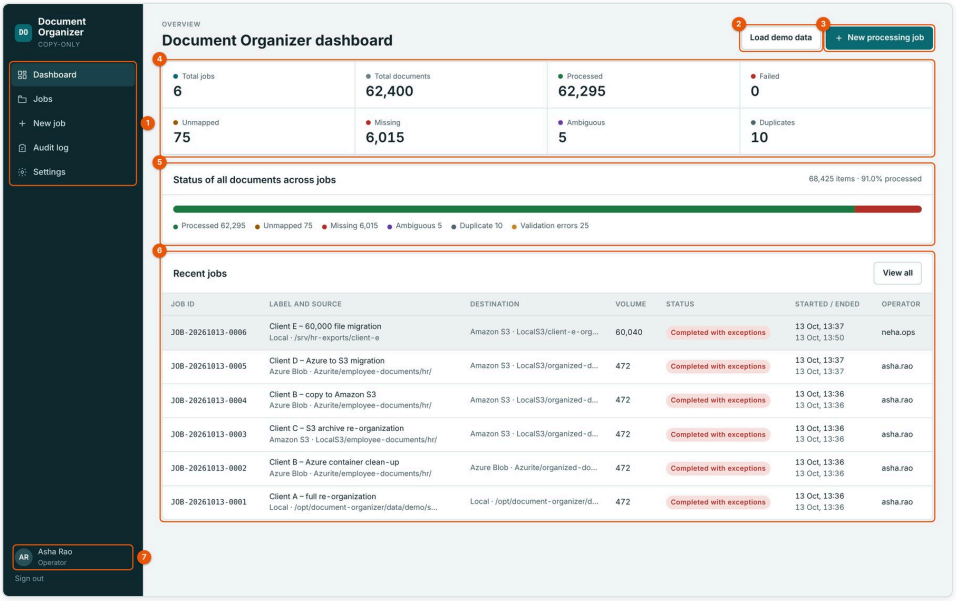


Figure 3: Dashboard

#	Item	Description
1	Navigation bar	Moves between Dashboard, Jobs, New job, Audit log and Settings.
2	**Load demo data**	Creates fictional practice files and a matching mapping file so you can try the full flow safely. See Appendix A. Shown only when demo data is enabled.
3	**New processing job**	Starts a new job (section 5). Operators and administrators only.
4	Totals	Total jobs, total documents found in sources, and totals per status across all jobs: Processed, Failed, Unmapped, Missing, Ambiguous and Duplicates.
5	Status of all documents	A bar showing the share of each status across all jobs, with counts in the legend.
6	Recent jobs	The latest jobs with job ID, label and source, destination, volume (files found), status, start and end time, and operator. Click a row to open the job.
7	Signed-in user	Your name and role. **Sign out** is just below.

Reading the totals

"Missing" counts mapping records whose file was not found in the source. A large number usually means the source prefix or folder was too narrow, not that files are lost. Check the job's exception list before reporting to the client.

4 Jobs list

Jobs lists every job, newest first. Click any row to open that job at its current step.

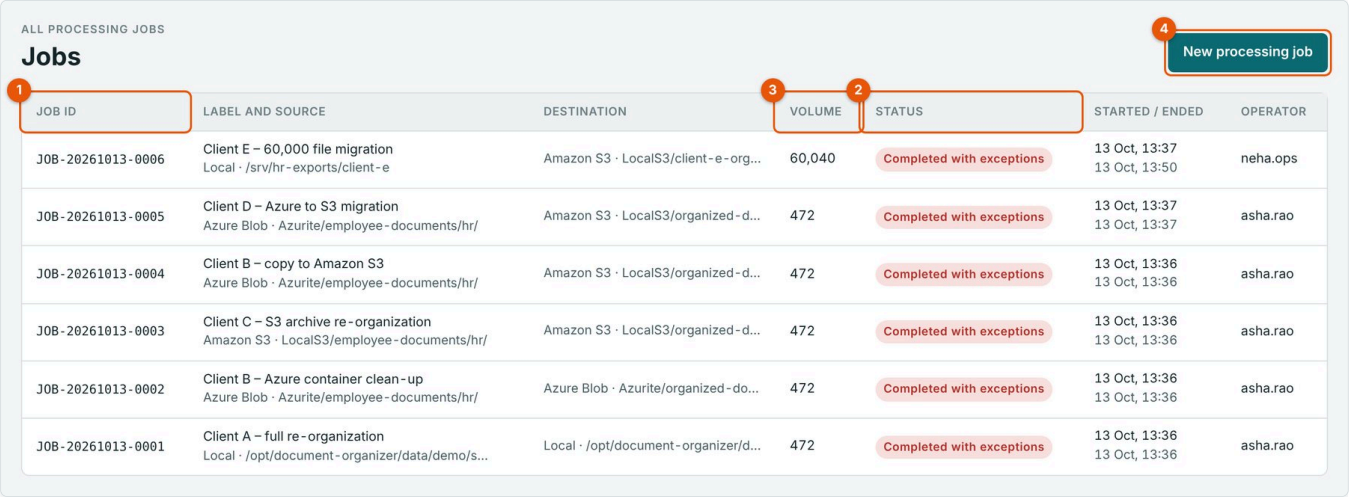


Figure 4: Jobs list

#	Item	Description
1	Job ID	A unique ID in the form `JOB-YYYYMMDD-NNNN`. Quote it when you ask for help.
2	Status	Where the job is. See the table below.
3	Volume	Number of files found in the source by the scan.
4	**New processing job**	Starts a new job.

Job status	Meaning	What you can do
Draft	Created; the mapping file and rules are being set up, or the rules were changed after a scan.	Continue the setup and scan.
Scanning	The source is being listed and matched.	Wait; you can leave the page.
Awaiting approval	Scan finished. Results are ready for review.	Review, resolve, download the validation report, approve.
Processing	Files are being copied.	Watch progress or pause.
Paused	Copying was paused by a user or interrupted by a restart.	Resume.
Completed	Every item finished with no exceptions.	Download the reconciliation report.
Completed with exceptions	Copying finished; some items need follow-up (for example unmapped or missing).	Download the reports; retry failed files if any.

5 Creating a job: source and destination

A job goes through eight steps, shown in the step bar at the top of the job page. This section covers step 1. Click **New job** (or **New processing job** on the dashboard).

5.1 Choose the source

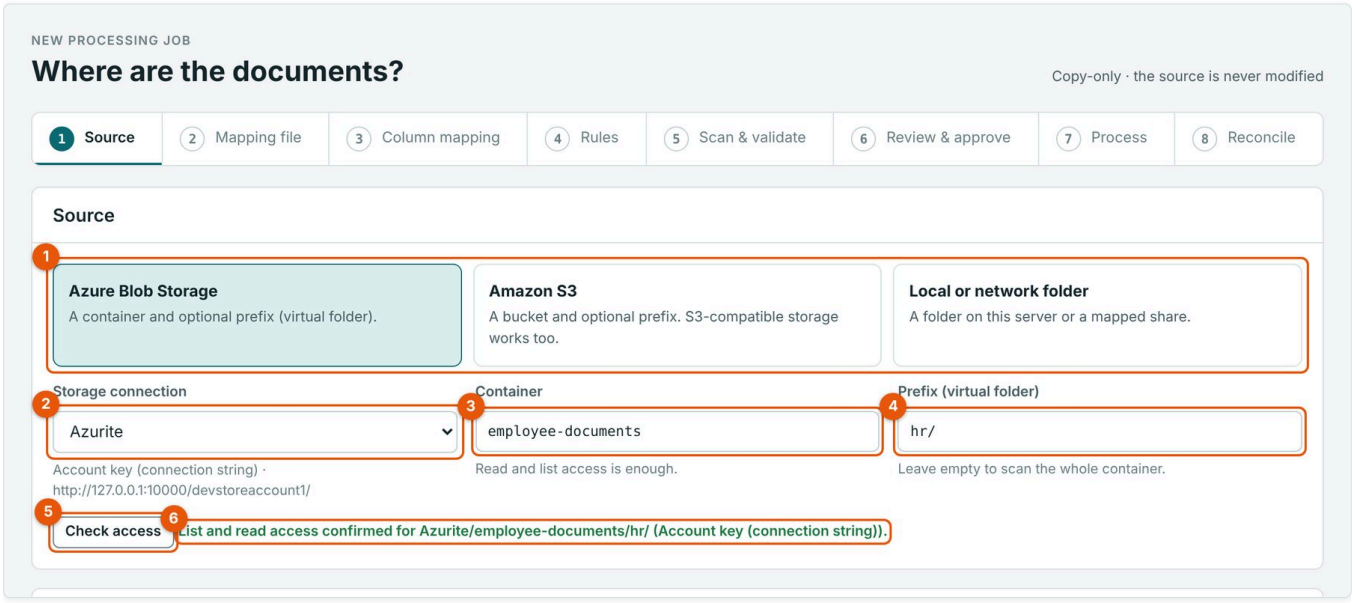


Figure 5: Step 1, source

#	Item	What to do
1	Storage type	Choose **Azure Blob Storage**, **Amazon S3** or **Local or network folder**. A type is greyed out if your administrator has turned it off or no connection of that type is configured.
2	**Storage connection**	Pick the storage account. Connections are set up by your administrator; only connections of the chosen type are listed. The line below shows how the application signs in and the account address.
3	**Container** (Azure) or **Bucket** (S3)	The container or bucket that holds the documents, for example `employee-documents`.
4	**Prefix (virtual folder)**	The folder inside the container or bucket, for example `hr/`. Leave it empty to scan everything.
5	**Check access**	Tests that the application can list and read the source.
6	Result	Green text confirms access. Red text explains the problem, for example a wrong name, a bucket in another region, or a missing permission.

For a folder source, type the full path instead, for example `D:\HRDocs\ClientA` or `/srv/hr-exports/client-a`. It is scanned including all sub-folders.

5.2 Choose the destination and name the job

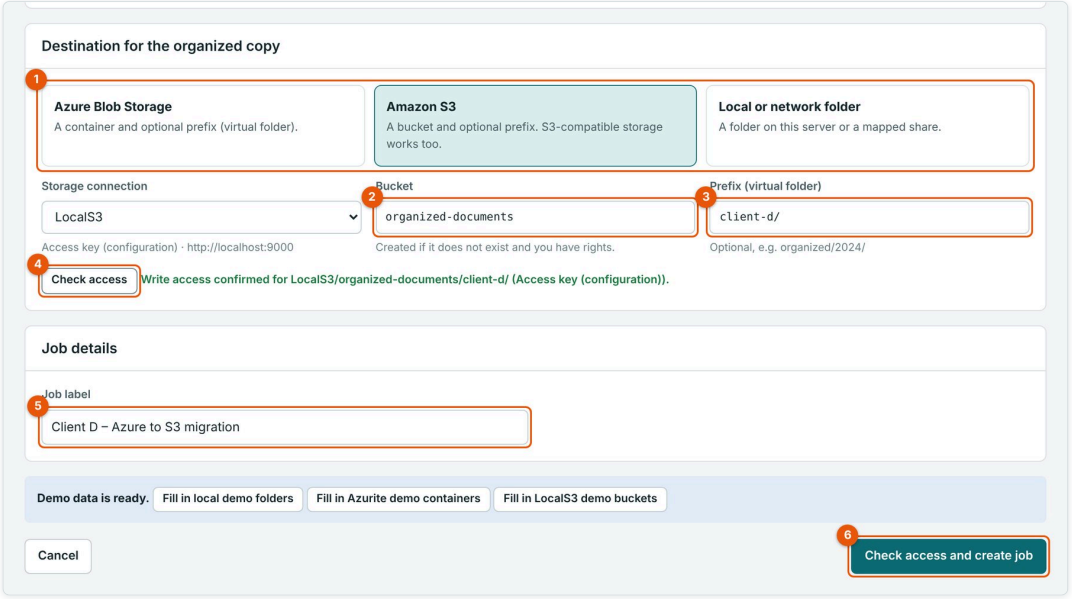


Figure 6: Step 1, destination and job details

#	Item	What to do
1	Storage type	Choose where the organized copy is written. It can be the same type as the source or different (for example Azure source, S3 destination).
2	**Container / Bucket**	The destination container or bucket, for example `organized-documents`. It is created if it does not exist and the connection has the rights.
3	**Prefix**	Optional folder inside the destination, for example `client-d/`.
4	**Check access**	Tests that the application can write to the destination. A small test file is written and removed.
5	**Job label**	A short name your team will recognise, for example "Client D – Azure to S3 migration".
6	**Check access and create job**	Checks both locations again and creates the job. You move to step 2.

> **The destination must be separate from the source**
>
> The application refuses a destination folder inside the source folder, or a destination prefix inside the source prefix in the same container or bucket. This prevents the organized copy from being picked up by a later scan.

5.3 Filling in the demo locations

After **Load demo data** has been used, a blue bar appears on this page.



Figure 7: Demo shortcuts on the New job page

#	Item	Description
1	**Fill in local demo folders**	Fills in the demo source and destination folders on the application server.
2	**Fill in … demo containers**	Fills in the demo Azure containers. Shown for each Azure connection marked for demo use that was running when the demo data was loaded (for example the Azurite emulator).
3	**Fill in … demo buckets**	The same for S3 connections marked for demo use (for example MinIO).

6 Uploading the mapping files

Step 2 asks for the mapping file: a spreadsheet with one row per document, usually the document log exported from the HR system as-is. You can upload **several mapping files** for one job, for example one export per module (HR documents, payroll, onboarding). Their records are combined in one scan, and every record keeps the name of the file it came from.

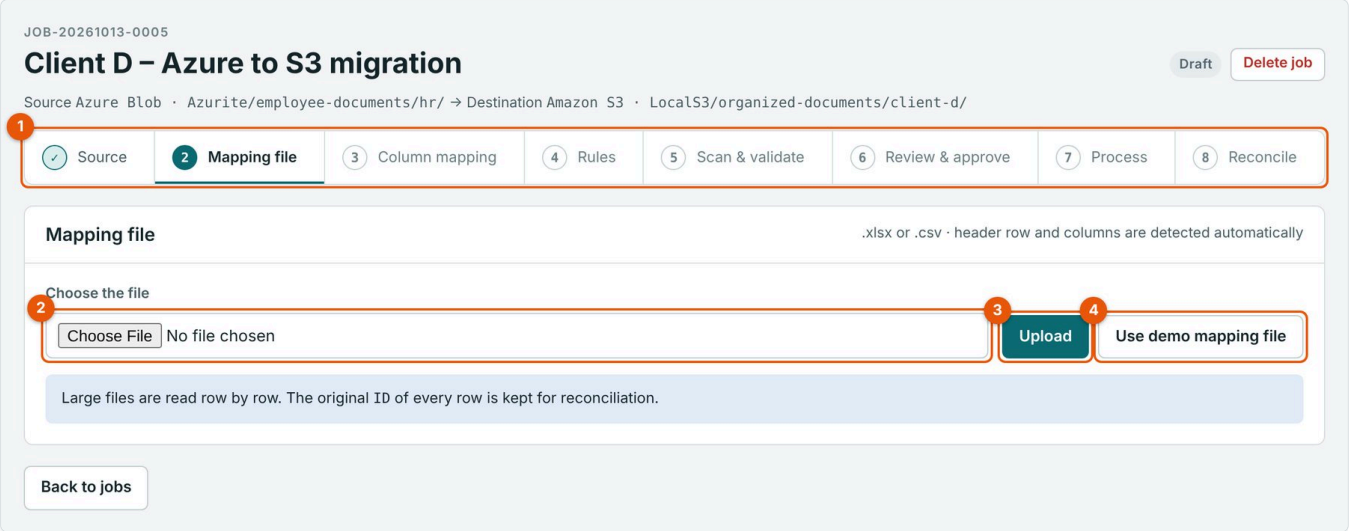


Figure 8: Step 2 before upload

#	Item	What to do
1	Step bar	Shows the eight steps. Completed steps have a tick. You can click steps 2 to 4 to go back while the job is not yet approved.
2	**Choose the file(s)**	Select one or more mapping files (`.xlsx` or `.csv`); hold Ctrl or Shift to select several. Up to 300 MB per upload.
3	**Upload**	Uploads and reads the files. A 176,000-row Excel file takes about 3 seconds.
4	**Use demo mapping file**	Uses the demo mapping file (practice only).

After the upload the page lists every mapping file:

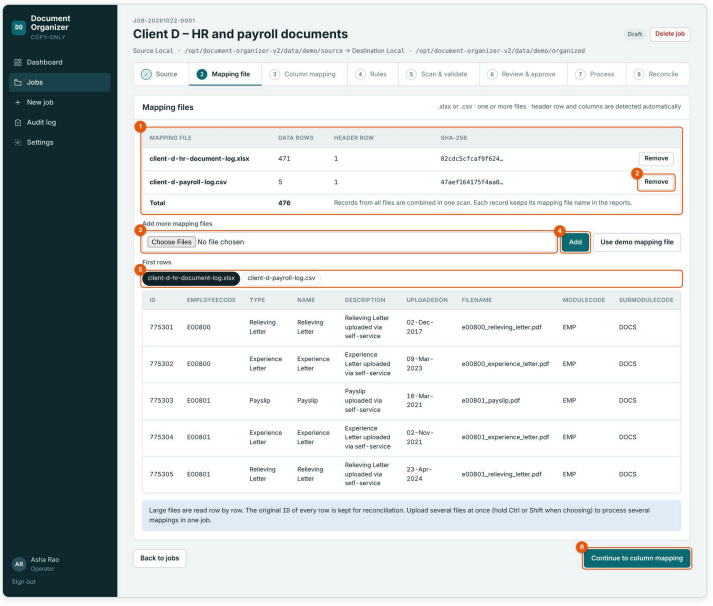


Figure 9: Step 2 with two mapping files

#	Item	Description
1	Mapping files	One row per file: name, number of data rows, the row where the column headers were found, and the file's SHA-256 fingerprint (recorded in the audit log and reports). With more than one file, a **Total** row adds them up.
2	**Remove**	Removes that mapping file from the job.
3	**Add more mapping files**	Choose more files to add to the job.
4	**Add**	Uploads and reads the chosen files. Adding or removing a file after a scan clears the scan results.
5	First rows	The first five data rows of the selected file. With several files, click a file name to see its rows.

#	Item	Description
6	**Continue to column mapping**	Moves to step 3.

No reformatting needed

The application finds the header row by itself, even if the export has title rows above it. In an Excel workbook it uses the first visible sheet that has a header row, so a cover or summary sheet in front of the data is fine. Dates in Excel date format and text dates such as `25-Nov-2021` are both understood. CSV files may use commas, semicolons or tabs, and may be saved as UTF-8 or by Excel on Windows (accented names such as `Résumé.pdf` are read correctly). Cells that only say `NULL`, `N/A`, `NA`, `None` or `-` (common in database exports) are treated as empty.

Upload checks

A file with the same content as one already added is refused (it would turn every record into a duplicate). Two different files with the same name are kept apart by adding a number, for example `export (2).csv`. If any file in an upload cannot be read, nothing from that upload is kept and the job's current mapping files are untouched.

Several file names in one cell

Some exports list several uploaded files for one document in a single File name cell, separated by commas or semicolons, for example `marksheet_1.jpg,marksheet_2.pdf`. Each name becomes its own entry with its own status: record `ROW-68` becomes `ROW-68/1` and `ROW-68/2`. A name that really contains a comma is kept whole when a file with exactly that name is in the source, even next to other names: `report, final.pdf,offer.pdf` gives `report, final.pdf` and `offer.pdf`. Folder paths in the cell work too. A cell that only holds placeholders, such as `NULL,NULL`, counts as blank. You can turn splitting off in step 4.

7 Mapping the columns

Step 3 links each column in the mapping file to the field the application needs. Columns are matched automatically by name; you only change a row when the file uses a different column name. With several mapping files, each file has its own column mapping, so files with different column names, column order, extra columns or title rows can be combined. If a file's columns are not recognised, map them by hand on its tab; the scan is refused until every file has its required columns, and the message names the file.

Files with different date columns

The year rule (step 4) applies to the whole job, but files can differ. With the rule **UploadedOn year**, a file that has no Uploaded on column but has a Year column uses its Year column; with **Dedicated Year column**, a file without a Year column uses its Uploaded on date. Step 3 shows a note on that file's tab when this applies.

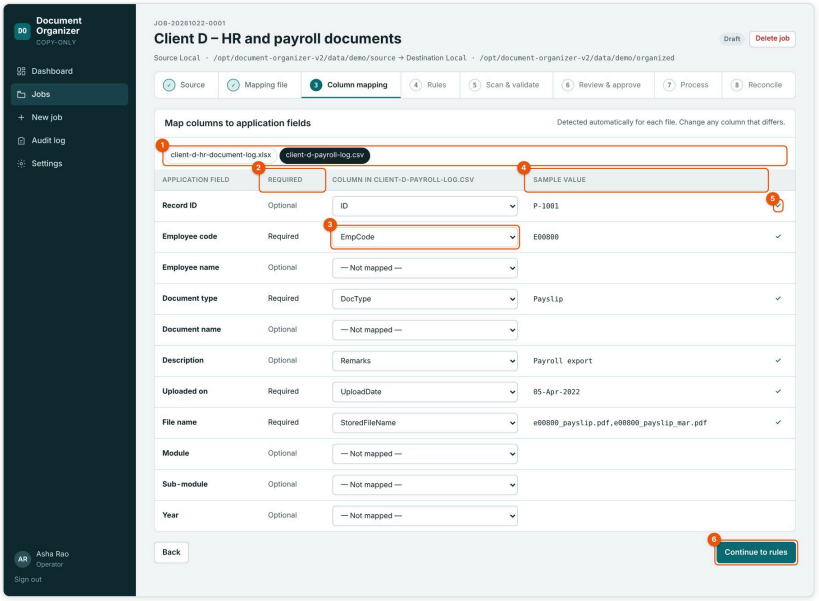


Figure 10: Step 3, column mapping for the second of two mapping files

#	Item	Description
1	Mapping file	Shown when the job has several mapping files. Click a file name to see and change its column mapping. A warning lists any file that still has a required column missing.
2	Required	Employee code, Document type and File name are always required. Uploaded on is required when the year comes from UploadedOn; Year is required when the year comes from a Year column.
3	Column in file	The column used for this field. Choose another column or **— Not mapped —**.
4	Sample value	The value in the first data row, to confirm the choice.
5	Check	A tick when the field is mapped; **Missing** in red when a required field is not mapped.
6	**Continue to rules**	Moves to step 4.

Application field	Typical column names	Used for
Record ID	`ID`, `RecordID`, `DocumentID`	Traceability in reports. If missing, the row number is used.
Employee code	`EmployeeCode`, `EmpCode`, `EmployeeID`	First folder level.
Employee name	`EmployeeName`, `EmpName`	Reports, optional template token.
Document type	`Type`, `DocumentType`, `DocType`, `Category`	Second folder level.
Document name	`Name`, `DocumentName`, `Title`	Reports, optional template token.
Description	`Description`, `Remarks`	Reports.
Uploaded on	`UploadedOn`, `UploadDate`, `CreatedOn`	Year (default rule).
File name	`FileName`, `PhysicalFileName`, `FilePath`, `StoredFileName`	Matching mapping records to files. A bare name or a relative path; several names separated by commas or semicolons are split (section 6).
Module / Sub-module	`ModuleCode` / `SubModuleCode`	Reports and optional template tokens.
Year	`Year`, `DocumentYear`, `FinancialYear`	Year (when the Year column rule is chosen).

Employee name, document name, module and sub-module are optional, but if the folder template uses one of them, its column must be mapped (the rules are refused otherwise), and a record whose value is blank, `NULL` or only dots becomes INVALID. That way no file lands one folder level too high.

8 Setting the rules

Step 4 decides how the year is worked out, how folders are named, how files are matched, what happens when a file matches several records or already exists at the destination, and how copies are verified.

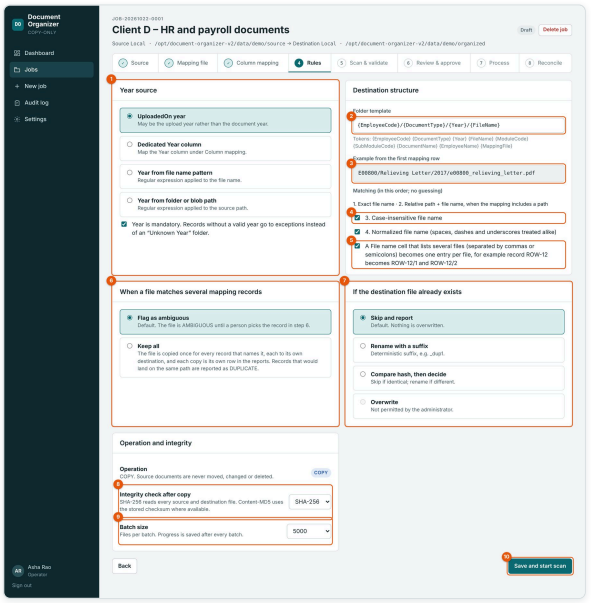


Figure 11: Step 4, rules

#	Item	Description
1	Year source	**UploadedOn year** (default), **Dedicated Year column**, **Year from file name pattern** or **Year from folder or blob path**. For the two pattern options, a **Pattern** box appears; the first 4-digit year found is used. Below it, **Year is mandatory** sends records without a valid year to exceptions instead of an "Unknown Year" folder.
2	**Folder template**	How destination paths are built. Default: `{EmployeeCode}/{DocumentType}/{Year}/{FileName}`. Available tokens are listed under the box. It must end with `{FileName}`.
3	Example	The destination path the template gives for the first mapping record. It updates as you type.
4	Matching options	Exact file name and path matching always apply. **Case-insensitive file name** and **Normalized file name** (spaces, dashes and underscores treated alike) are on by default.
5	Several file names in one cell	On by default. A File name cell listing several files, separated by commas or semicolons, becomes one entry per file (section 6).
6	When a file matches several mapping records	**Flag as ambiguous** (default): the file is AMBIGUOUS until a person picks the record in step 6. **Keep all**: the file is copied once for every record that names it, each to its own destination (for example `E1/Offer/2024/x.pdf` and `E2/PAN/2023/x.pdf`), and each copy is its own row in both reports. If two of those records would give the same destination path, one copy is kept and the other is DUPLICATE.
7	If the destination file already exists	**Skip and report** (default), **Rename with a suffix** (`_dup1`, `_dup2`), **Compare hash, then decide** (skip if identical, rename if different) or **Overwrite** (only when an administrator allows it).
8	**Integrity check after copy**	**SHA-256** (strongest; reads every file twice), **Content-MD5** (uses the checksum the storage already holds where it can, which is cheaper on Azure and for single-part S3 objects) or **None**.
9	**Batch size**	Files per batch (1,000, 5,000 or 10,000). Progress is saved after every batch.
10	**Save and start scan**	Saves the rules and starts step 5.

For cloud destinations an extra switch, **Keep object metadata**, keeps the source object's metadata and adds the job ID, record ID and employee code to each copied file.

The token `{MappingFile}` puts each document under the name of the mapping file it came from (without the extension), for example `{MappingFile}/{EmployeeCode}/{DocumentType}/{FileName}` gives `client-d-payroll-log/E00801/Payslip/e00801_payslip.pdf`. Use it to keep the documents of each mapping file apart.

> **Changing rules later**
>
> You can come back to steps 2 to 4 until the job is approved. Saving changed rules clears the scan results, and you scan again. After approval the rules are locked, so each job's report always matches the rules it ran with.

9 Scanning and validating

Step 5 runs by itself. The application reads the mapping file, lists every file in the source, matches files to mapping records and checks each one. You can leave the page; the scan continues in the background.

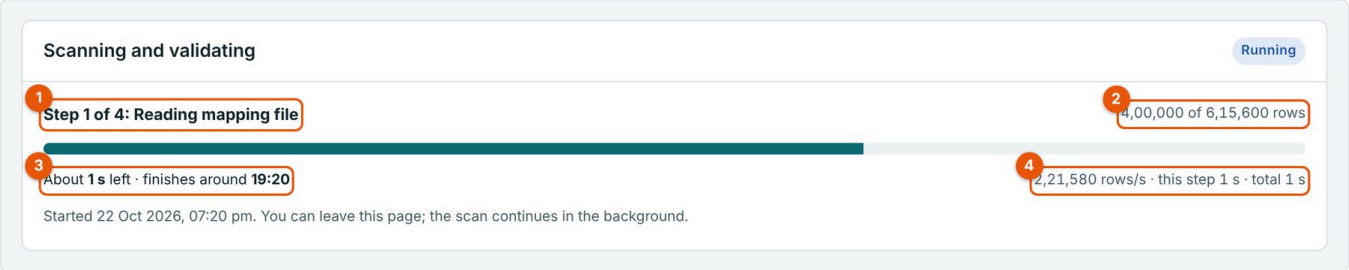


Figure 12: Step 5, scan in progress

#	Item	Description
1	Current step	Step 1 of 4 *Reading mapping file*, 2 *Listing source documents*, 3 *Matching and validating*, 4 *Saving results*.
2	Progress	Rows, files or items handled so far in this step, out of the total when it is known.
3	Time left	*About 2 min left · finishes around 19:20* for the current step, worked out from the speed of the last two minutes. While the source is being listed the total is not known yet, so the number of files found is shown instead.
4	Speed and time taken	Rows or files per second, how long this step has run, and the total time of the scan so far.

In testing, a 66,000-row mapping file with 60,000 files was scanned and validated in about 3 seconds from a local folder. Cloud sources take longer because files are listed in pages (5,000 per page on Azure, 1,000 per page on S3).

What is checked: blank employee code, type or file name (including placeholders such as `NULL`); a missing or unreadable year; duplicate mapping records; the same file found twice; two files that would land on the same destination path; empty (0-byte) files; files in an archive tier (Azure Archive, S3 Glacier Flexible Retrieval or Deep Archive); and mapping records whose file is not in the source.

10 Reviewing and approving

When the scan finishes the job shows **Awaiting approval** and step 6 opens. Nothing is copied until you approve.

10.1 Results summary

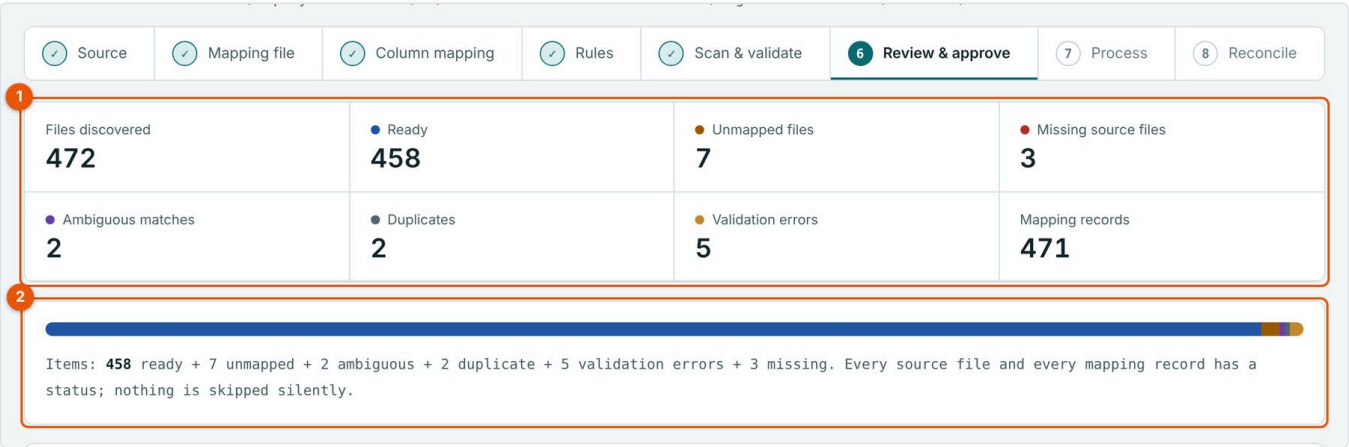


Figure 13: Step 6, results summary

#	Item	Description
1	Counts	Files discovered, Ready, Unmapped files, Missing source files, Ambiguous matches, Duplicates, Validation errors and Mapping records. Click a coloured tile to filter the table below to that status.
2	Check bar and line	The share of each status, and a line showing that every file and record is counted once.

10.2 Exceptions and results

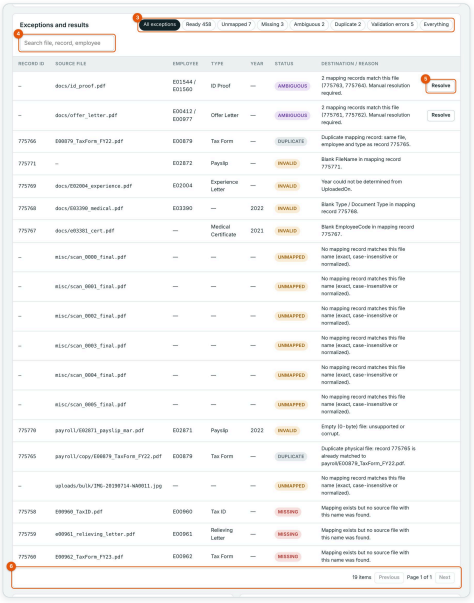


Figure 14: Step 6, exceptions table

#	Item	Description
3	Filters	**All exceptions** (default), one tab per status with its count, **Excluded**, or **Everything**. Tick boxes and bulk buttons for excluding are described in section 10.3.
4	Search	Finds items by file name, record ID, employee code, destination path or reason.
5	**Resolve**	Shown on AMBIGUOUS items. Opens the resolve panel (section 10.4).
6	Pages	50 items per page; use **Previous** and **Next**.

The **Destination / reason** column shows the proposed destination path for READY items, and the reason for every exception in plain words.

10.3 Excluding exceptions in bulk

Some exceptions need no follow-up: blank file names the client has confirmed, scanner test pages, files that belong to another module. **Exclude** them to take them out of the exception lists, with a note saying why. Excluded items:

- are never copied and no longer count as exceptions (a finished job with only processed and excluded items is **Completed**);
- keep their original status and reason, and stay in both reports with status **EXCLUDED**, who excluded them, when and the note;
- are excluded again automatically after a new scan, as long as the same file or record still has the same status;
- can be restored at any time from the **Excluded** tab.

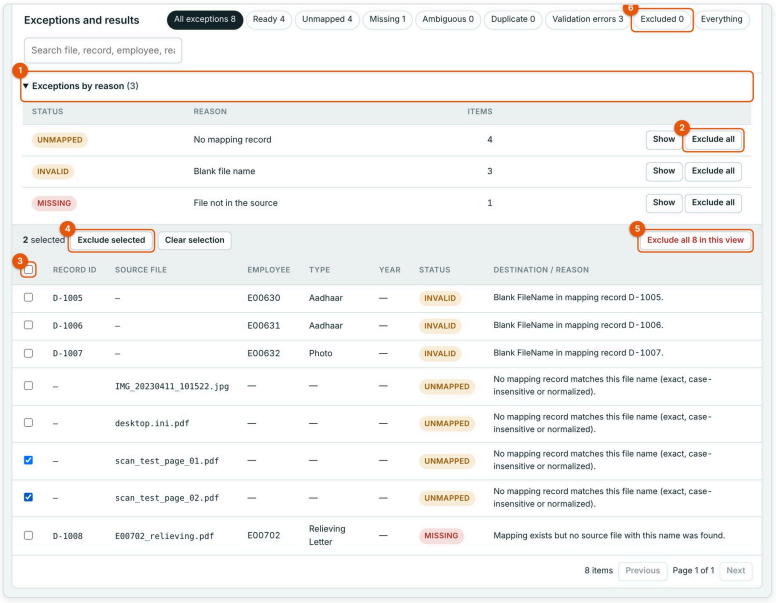


Figure 14a: Step 6, choosing exceptions to exclude

#	Item	Description
1	**Exceptions by reason**	Click to open the list of reasons with their status and number of items.
2	**Exclude all** (per reason)	Excludes every item with that reason. **Show** next to it filters the table to that reason.
3	Select page	Ticks every exception on this page. Each row also has its own tick box.
4	**Exclude selected**	Excludes the ticked rows. The count on the left shows how many are ticked.
5	**Exclude all N in this view**	Excludes everything the current tab and search show, on every page.
6	**Excluded** tab	Lists excluded items and lets you restore them.

There are three ways to choose what to exclude. All of them work on every exception type (unmapped, missing, ambiguous, duplicate, validation errors and, after copying, failed):

Way	How
Ticked rows	Tick the box on each row (the box in the header ticks the whole page; ticks are kept when you change page), then **Exclude selected**.
Everything in the view	Choose a tab and, if you like, search; then **Exclude all N in this view**.
A whole reason	Open **Exceptions by reason** above the table. Each reason (for example *Blank file name*, *No mapping record*, *Year not found*) is listed with its status and count. **Show** filters the table to that reason; **Exclude all** excludes it at once.

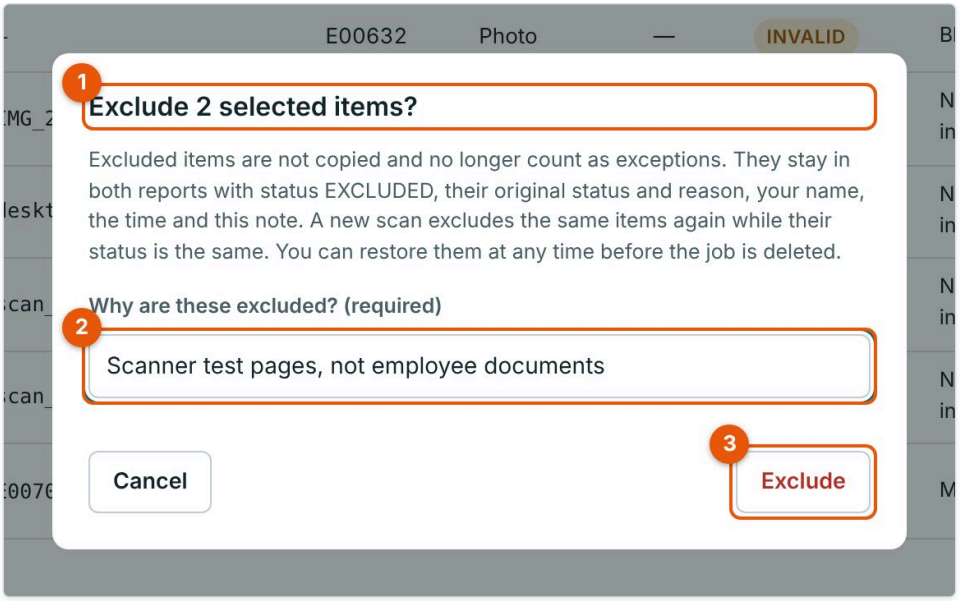


Figure 14b: The note is required

#	Item	Description
1	Title	Says exactly what will be excluded: the number of ticked items, or all items of a view or reason.
2	Note	Why these items need no follow-up. Kept on each item, in both reports and in the audit log.
3	**Exclude**	Confirms. **Cancel** changes nothing.

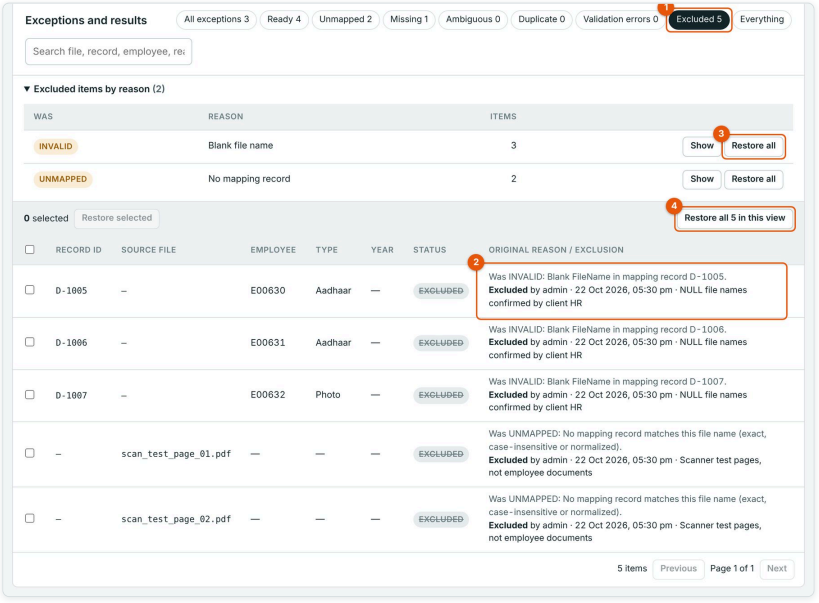


Figure 14c: The Excluded tab

#	Item	Description
1	**Excluded** tab	Every excluded item, with the count.
2	Original reason and exclusion	The status and reason the item had, then who excluded it, when, and the note.
3	**Restore all** (per reason)	Puts every item of that reason back to its original status.
4	**Restore all N in this view**	Restores everything shown. Tick rows and use **Restore selected** for single items.

A note is required each time and is recorded in the audit log. On the **Excluded** tab the same three ways restore items (**Restore selected**, **Restore all N in this view**, or **Restore all** for a reason); restoring needs no note. Operators and administrators can exclude and restore while the job is awaiting approval, paused or completed; viewers can see excluded items.

> Excluding is not a fix. If an exception means the mapping file is wrong, correct the mapping file and scan again instead.

10.4 Resolving an ambiguous file

A file is AMBIGUOUS when more than one mapping record names it, for example `id_proof.pdf` for two different employees. The application will not choose; a person must.

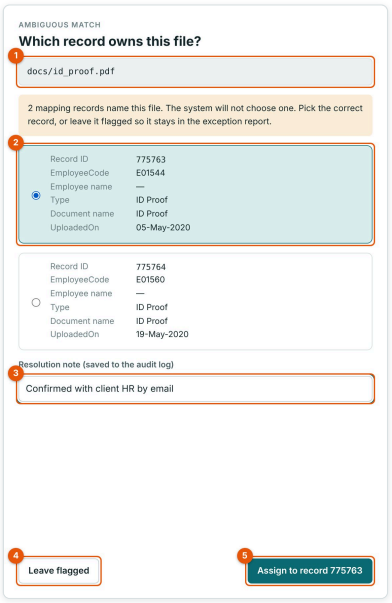


Figure 15: Resolve panel

#	Item	What to do
1	File	The file that needs a decision, with its location.
2	Candidate records	Each mapping record that names this file, with record ID, employee code and name, type, document name and upload date. Select the correct one.
3	**Resolution note**	Say how you decided, for example "Confirmed with client HR by email". Saved in the audit log.
4	**Leave flagged**	Closes the panel without a decision. The file stays AMBIGUOUS and is not copied.
5	**Assign to record ...**	Assigns the file. It becomes READY, or DUPLICATE if that record or destination is already taken.

10.5 Checking the output and approving

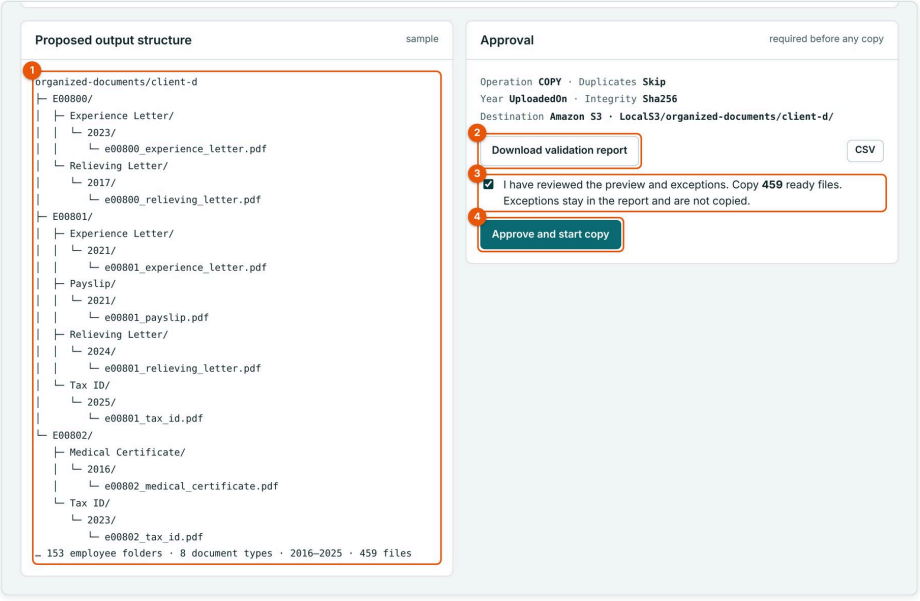


Figure 16: Step 6, proposed structure and approval

#	Item	Description
1	Proposed output structure	A sample of the destination folders for the first employees, and totals: employee folders, document types, year range and files.
2	**Download validation report**	Excel (or **CSV**) list of every item with its status and reason. Use it to agree exceptions with the client before copying.
3	Confirmation	Tick to confirm you reviewed the preview and exceptions.
4	**Approve and start copy**	Opens the confirmation below.

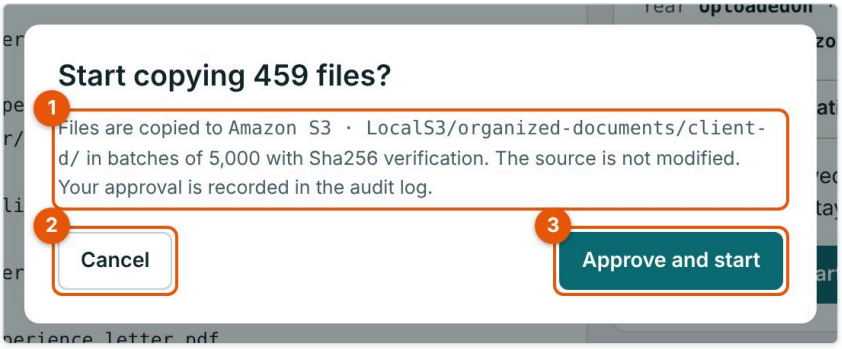


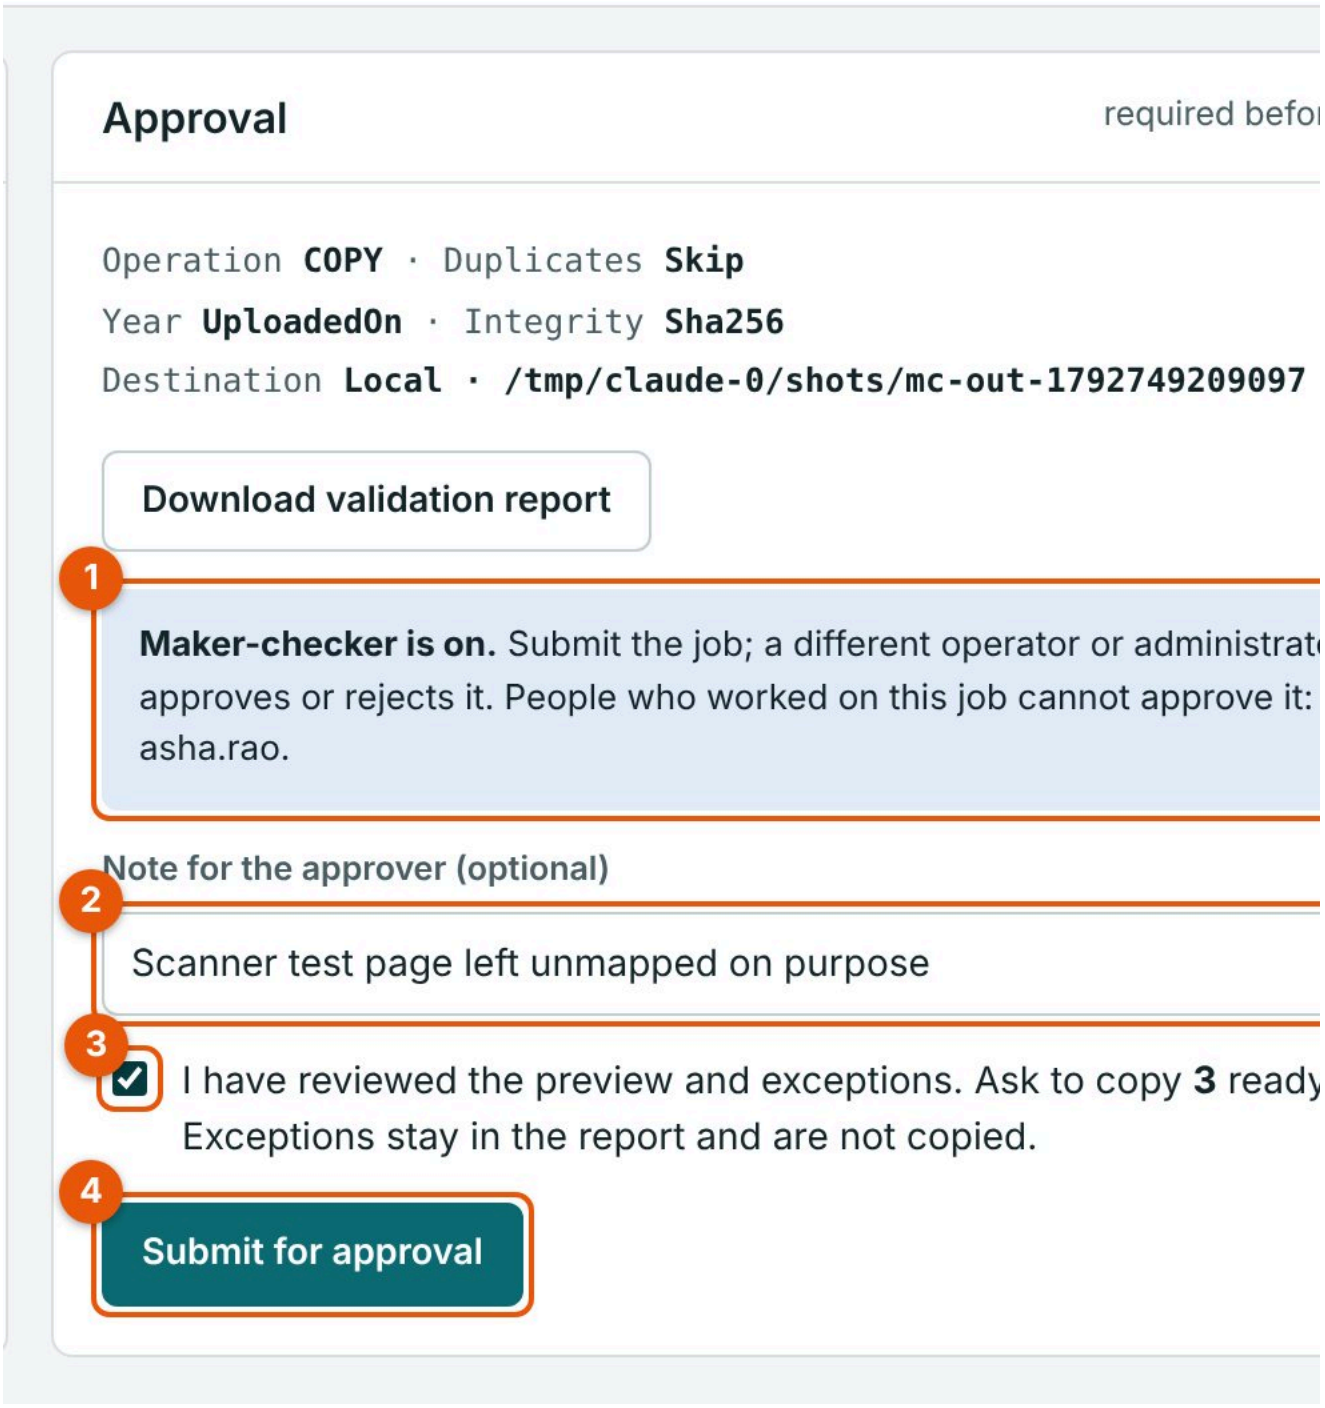
Figure 17: Approval confirmation

#	Item	Description
1	Summary	Number of files, the destination, batch size and integrity check.
2	**Cancel**	Goes back without copying.
3	**Approve and start**	Records your approval (name and time) in the audit log and starts copying.

10.6 Maker-checker (second approver)

When an administrator turns on **Maker-checker** in Settings, no one can approve a job they worked on. This is the four-eyes control many audit teams require.

- **Makers.** Everyone who creates the job, uploads or removes mapping files, saves rules, scans, resolves ambiguous files, or excludes or restores items is recorded as a maker of that job.



#	Item	Description
1	Maker-checker notice	Shown when the setting is on. Lists the people who worked on the job and so cannot approve it.
2	Note for the approver	Optional. Shown to the checker with the job.
3	Confirmation	Required before submitting.
4	**Submit for approval**	Locks the job and sends it to a checker.

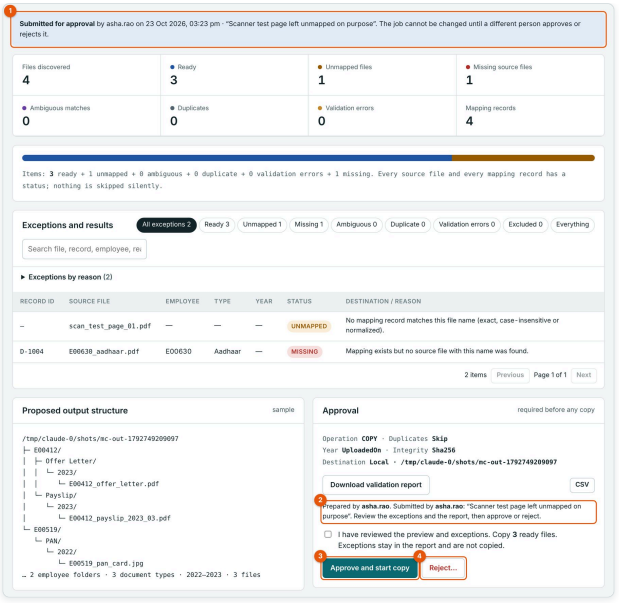


Figure 14e: What the checker sees

#	Item	Description
1	Submitted banner	Who submitted the job, when, and the note. The job cannot be changed while it waits.
2	Prepared and submitted by	The makers and the submitter.
3	**Approve and start copy**	Starts copying after the confirmation. Refused if the job changed since the page was opened (reload and review again).
4	**Reject...**	Sends the job back to the makers with a reason (required).

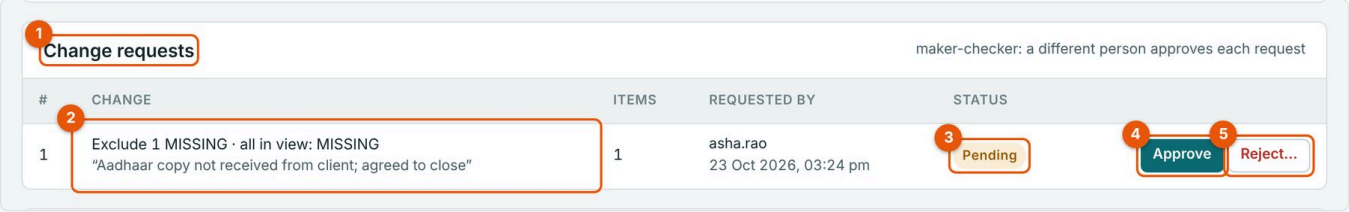


Figure 14f: Change requests after approval

#	Item	Description
1	Change requests	On the reconciliation step when maker-checker is on.
2	Change	What would change (exclude or restore, items and filter) and the requester's note.
3	Status	Pending, Approved, Rejected or Withdrawn, with who decided.
4	**Approve**	For anyone except the requester: applies the change.
5	**Reject...**	With a reason. The requester sees **Withdraw** instead of these buttons.

- **Submit.** When the scan is reviewed, a maker ticks the confirmation and clicks **Submit for approval** (with an optional note for the approver). The job shows **Submitted for approval** and is locked: nobody can change mapping files, rules, exclusions or resolutions until it is approved, rejected or withdrawn. The person who submitted it, or an administrator, can **Withdraw** it to make changes.
- **Approve or reject.** Any operator or administrator who is not one of the job's makers sees **Approve and start copy** and **Reject...**. Rejecting needs a reason, which the makers see on the job. People who worked on the job see who can approve it instead of the buttons.
- **After approval.** Excluding or restoring items on an approved job sends a **change request**. Nothing changes until a different person approves it under **Change requests** on the reconciliation step. The requester can withdraw it; anyone else can approve it or reject it with a reason. An approved exclusion keeps the requester as *Excluded by* and adds *(approved by ...)* to the note.
- **Records.** Every submission, withdrawal, rejection, approval and change request is in the audit log. The report summary shows who prepared, submitted and approved the job.

With maker-checker off (the default), the person who prepares a job approves it directly, as before. Viewers can never approve.

> Someone who changes a job, even to fix one exclusion, becomes one of its makers and can no longer approve it. In a team of two, the person who did not touch the job approves it.

11 Processing (copying)

Step 7 copies READY files in batches. When source and destination are in the same Azure storage account, or use the same S3 connection, files are copied inside the storage service and never pass through the application server. Otherwise (for example Azure to S3) the application streams each file from the source to the destination.

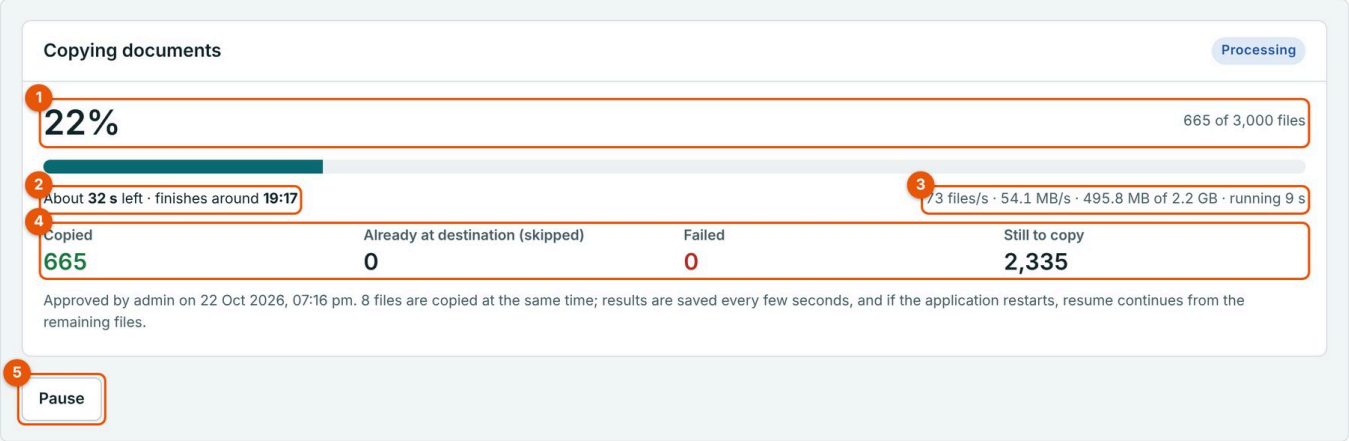


Figure 18: Step 7, copying in progress

#	Item	Description
1	Progress	Percentage and number of files done out of the total.
2	Time left	*About 1 h 20 min left · finishes around 21:40*. Worked out from the speed of the last two minutes and the size of the files still to copy, so it adjusts if the pace changes (for example many large scans near the end).
3	Speed	Files per second, megabytes per second, data copied so far out of the total, and how long the copy has run.
4	Counts	**Copied** (copied and verified), **Already at destination (skipped)** (the file existed there and the duplicate rule is Skip), **Failed**, and **Still to copy**. These count only what the copy step did; duplicates found during the scan are in the review and reconciliation screens.
5	**Pause**	Stops after the files in progress finish.

Several files are copied at the same time (8 by default) and results are saved every two seconds, so the counts move steadily.

Copying faster

Speed depends mostly on the drives or storage involved, not on the application. Folder-to-folder copies with an integrity check read each source file once (it is checked while it is copied) and read the copy back once to verify it. To go faster:

- Ask an administrator to change **Parallel copies** in Settings: 16 to 32 for SSD, NVMe or network storage; 2 to 4 for a single external USB hard disk (more at once makes a hard disk slower).
- Copy to a different physical drive from the source when you can.
- Exclude the source and destination folders from real-time antivirus scanning for the duration of the copy, if your IT policy allows it; scanning every new file can slow a copy several times over.
- Use **Content-MD5** instead of SHA-256 only for cloud storage that already holds checksums; for folders both read the files.

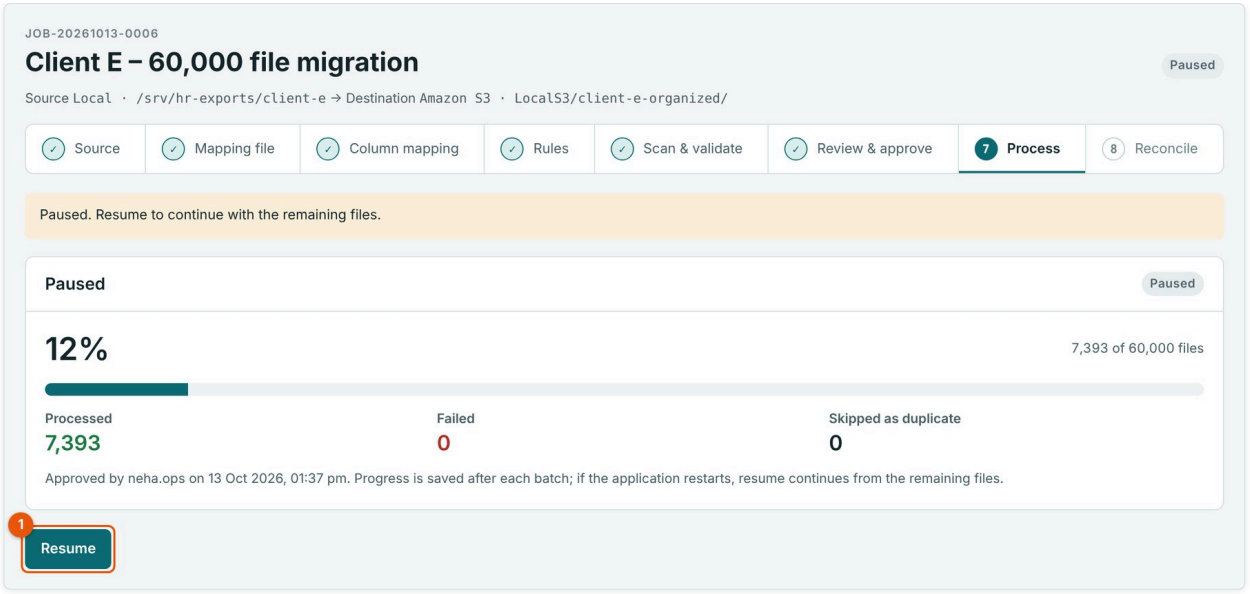


Figure 19: A paused job

#	Item	Description
1	**Resume**	Continues with the remaining files. Files already copied are not copied again.

If the application or server restarts

A job that was copying is shown as **Paused** with a message. Open it and click **Resume**. Files that were copied just before the restart are recognised (by comparing their contents) and recorded as processed, not copied twice.

A file is not copied, and is marked FAILED, if it changed or was removed after the scan. This protects against copying a different version from the one you approved.

12 Reconciliation and reports

When copying finishes the job moves to step 8.

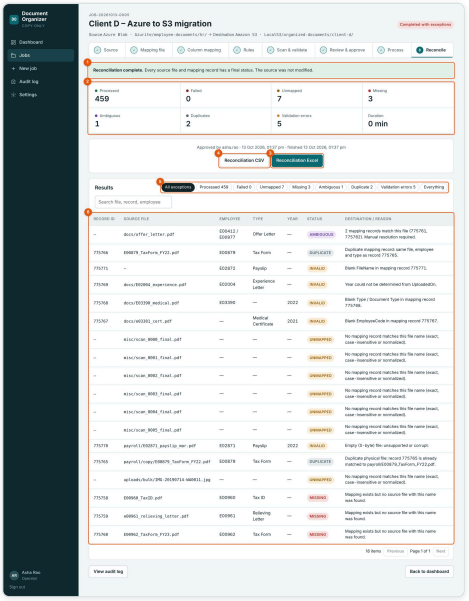


Figure 20: Step 8, reconciliation

#	Item	Description
1	Result	Green when every item has a final status; amber when some files failed and can be retried.
2	Final counts	Processed, Failed, Unmapped, Missing, Ambiguous, Duplicates, Validation errors, Excluded (when any) and Duration. Click a tile to filter the results.
3	**Reconciliation Excel**	Downloads the full report with a Summary sheet.
4	**Reconciliation CSV**	The same data as CSV.
5	Filters	Show exceptions, one status, or everything.
6	Results table	Every item with its destination (processed) or reason (exceptions).

Exceptions can still be excluded or restored here, as in section 10.3; when no exception is left the job becomes **Completed**.

When files failed, a **Retry N failed** button appears. Retrying copies only the failed files, using the current version of each file in the source. Files already processed are left alone.

12.1 What the reports contain

Report	Contents
Validation report (before approval)	Summary sheet (job, source, destination, every mapping file with its row count and fingerprint, rules, counts) and one row per item: item ID, status, reason, record ID, mapping file, employee code and name, document type and name, description, uploaded on, year, source file and path, proposed destination, size, access tier, and for excluded items the original status, who excluded them, when and the note.
Reconciliation report (after copying)	Summary sheet plus one row per item: job ID, record ID, mapping file, employee code and name, document type and name, description, uploaded on, source file and path, destination path, year, status, exception reason, source hash, destination hash, integrity method, source ETag, processed time (UTC) and processed by, then the same four exclusion columns.

13 Audit log

The audit log records every sign-in, job creation, connection check, mapping upload, configuration change, scan, resolution, approval, run, pause, resume, retry, report download, settings change and user change.

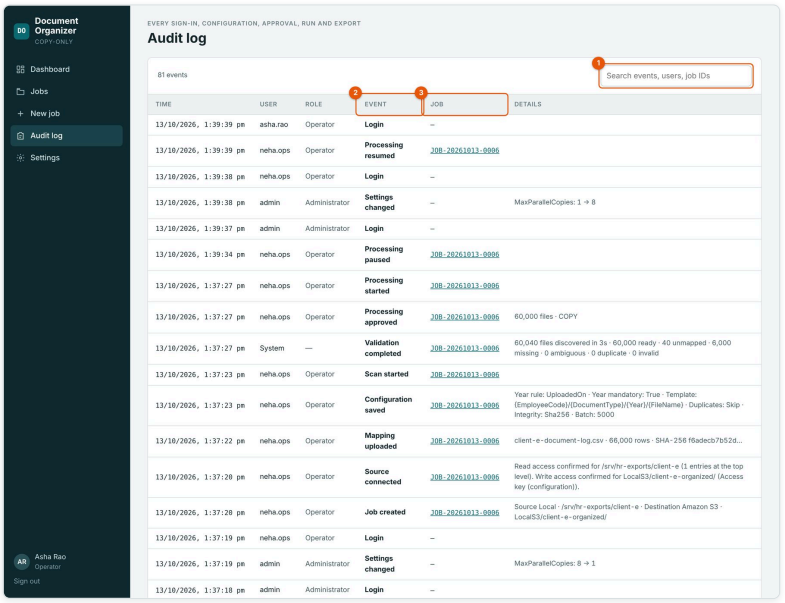


Figure 21: Audit log

#	Item	Description
1	Search	Finds events by event name, user, job ID or details.
2	Event	What happened, for example Processing approved or Ambiguity resolved.
3	Job	The job the event belongs to. Click it to open the job.

The log is shown newest first, 100 events per page. It cannot be edited from the application and is kept for 7 years by default.

14 Settings

Settings shows the application settings to everyone; only administrators can change them. Every change is recorded in the audit log.

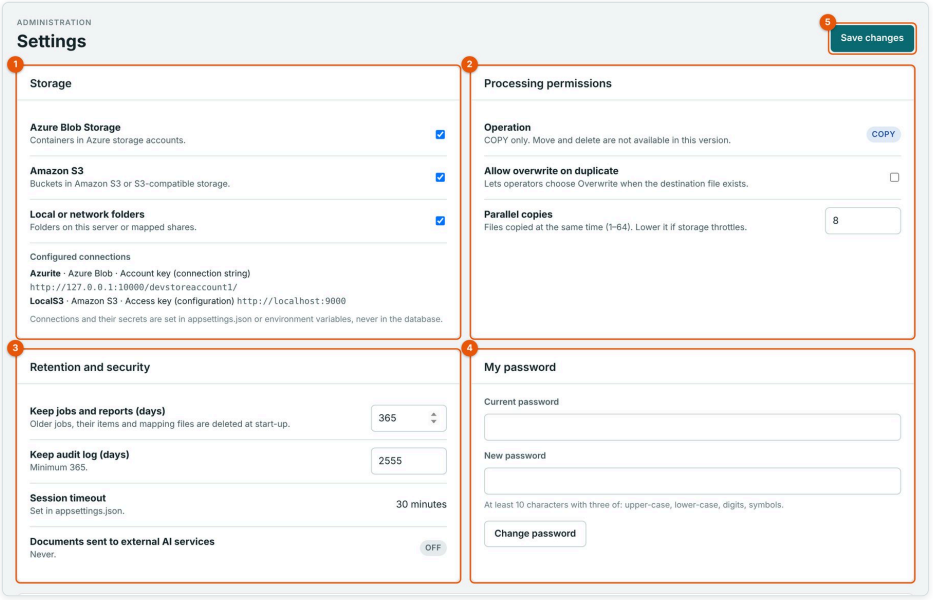


Figure 22: Settings

#	Item	Description
1	Storage	Turns Azure Blob Storage, Amazon S3 and local folders on or off, and lists the configured storage connections, their type and how each signs in.
2	Processing permissions	The COPY-only operation, **Allow overwrite on duplicate**, **Maker-checker (second approver)** (section 10.6), and **Parallel copies** (how many files are copied at the same time, 1 to 64).
3	Retention and security	How long jobs, reports and the audit log are kept, and the session timeout.
4	My password	Change your own password: current password, new password, **Change password**.
5	**Save changes**	Saves the settings (administrators).

14.1 Users and roles (administrators)

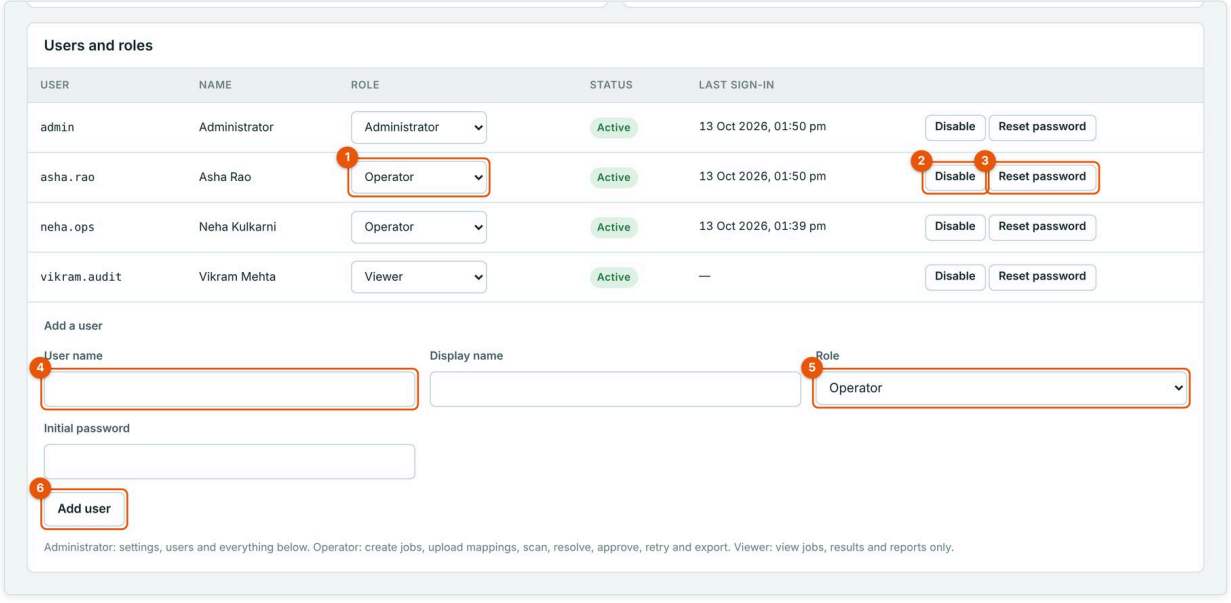


Figure 23: Users and roles

#	Item	What to do
1	Role	Change a user's role. It takes effect on their next click.
2	**Disable** / **Enable**	Blocks or restores a user's access. Their history stays in the audit log.
3	**Reset password**	Sets a new password for the user. Share it securely.

#	Item	What to do
4	**User name** (and **Display name**)	For a new user: 3 to 64 letters, digits, dot, dash, underscore or @.
5	**Role**	Operator, Viewer or Administrator.
6	**Add user**	Creates the user with the **Initial password**.

Passwords need at least 10 characters with three of: upper-case, lower-case, digits and symbols.

15 Status reference

Every source file and every mapping record ends with exactly one status. Only READY files are copied.

Status	Meaning	What to do
READY	The file matched one record and passed every check.	Nothing. It is copied after approval.
PROCESSED	Copied and verified.	Nothing.
UNMAPPED	A file in the source has no matching mapping record.	Check with the client or the data team whether the record is missing from the mapping file.
MISSING	A mapping record has no matching file in the source. A record that repeats another missing record is also MISSING, and its reason names the first record.	Check whether the file is in another container, bucket, prefix or folder.
AMBIGUOUS	More than one mapping record names the same file (with the default rule **Flag as ambiguous**).	Resolve it (section 10.4), leave it flagged, or scan with the rule **Keep all** to copy it for every record.
DUPLICATE	A repeated mapping record whose file was found (the first record is copied), a second copy of the same file, two files that would land on the same path, or a file already at the destination.	Read the reason. Usually no action; the first copy is kept.
INVALID	A validation error: blank employee code, type or file name (including `NULL`); a blank or dots-only value for a field the folder template uses; no usable year; an empty file; or a file in an archive tier.	Fix the mapping file or the file, upload the corrected mapping file and scan again.
FAILED	The copy did not succeed, for example storage was busy or the file changed after the scan.	Click **Retry failed**.
EXCLUDED	An exception that a person set aside with a note (section 10.3). Not copied; not counted as an exception. The original status and reason are kept.	Nothing, or **Restore** it from the Excluded tab.

16 Glossary

Term	Meaning
Azure Blob Storage	Microsoft's cloud file storage. Files live in containers inside a storage account.
Amazon S3	Amazon Web Services' cloud file storage. Files (objects) live in buckets. Many other services (MinIO, Wasabi, Cloudflare R2) offer the same interface.
Azurite	Microsoft's free Azure Storage emulator, used for testing on a laptop.
MinIO	A free S3-compatible storage server, used here as a local S3 emulator for testing.
Content-MD5	A checksum stored with many files; used as a cheaper integrity check.
ETag	A version marker of a file. A changed ETag means the file changed after the scan.
Archive tier	Low-cost storage that must be restored before files can be read: Azure Archive tier, S3 Glacier Flexible Retrieval and S3 Glacier Deep Archive. Not restored automatically.
Server-side copy	A copy made inside the storage service, without the content passing through the application.
SHA-256	A fingerprint of a file's contents; identical files have identical SHA-256 values.

17 Appendix A: practise with demo data

The demo creates about 470 fictional files and a matching mapping file that produce every status. It does not touch real data.

1. On the dashboard, click **Load demo data** (Figure 3, marker 2).
2. On the **New job** page, click **Fill in local demo folders** (or one of the demo container or bucket buttons if an emulator is running).
3. Click **Check access and create job**.
4. On step 2, click **Use demo mapping file**, then **Continue to column mapping** and **Continue to rules**.
5. Keep the default rules and click **Save and start scan**.
6. On step 6 you should see 472 files discovered, 471 mapping records, 458 ready, 7 unmapped, 3 missing, 2 ambiguous, 2 duplicate and 5 validation errors.
7. Resolve one ambiguous file (459 ready), tick the confirmation, approve, and download the reconciliation report.

Loading demo data again replaces the earlier demo files. Turn demo data off on production installations (Installation and Administration Guide, section 4.1).